

Neuro-Dynamic Programming

This repository is a collection of reading notes for the book [Neuro-Dynamic Programming](#) by Dimitri Bertsekas and John Tsitsiklis. Notes are provided as handwritten pdf files on a per-chapter basis. Following is the detailed outline of this repository-

Chapter-1: Introduction

Section	Description	Notes
1.1	Cost-to--go Approximations in Dynamic Programming	here
1.2	Approximation Architectures	here

Chapter-2: Dynamic Programming

Section	Description	Notes
2.1	Introduction	here
2.2	Stochastic Shortest Path Problems	here
2.3	Discounted Problems	here
2.4	Problem Formulation and Examples	here

Chapter-3: Neural Network Architectures and Training

Section	Description	Notes
3.1	Architectures for Approximation	here
3.2	Neural Network Training	here

Chapter-4: Stochastic Iterative Algorithms

Section	Description	Notes
4.1	The Basic Model	here
4.2	Convergence based on Smooth Potential Function	here
4.3	Convergence under Contraction or Monotonicity	here
4.4	The ODE Approach	here

Chapter-5: Simulation Methods for Lookup Table Representations

Section	Description	Notes
5.1	Some Aspects of Monte Carlo Simulation	here
5.2	Policy Evaluation by Monte Carlo Simulation	here

Section	Description	Notes
5.3	Temporal Difference Methods	here
5.4	Optimistic Policy Iteration	here
5.5	Simulation-Based Policy Iteration	here
5.6	Q-Learning	here

Chapter-6: Approximate Dynamic Programming with Cost-to-Go Function Approximation

Section	Description	Notes
6.1	Generic Issue- From Parameters to Policies	here
6.2	Approximate Policy Iteration	here
6.3	Approximate Policy Evaluation using TD(lambda)	here